

# Deep Learning Networks for Target Recognition



by ROAN SONG

prepared for A. K. MISHRA

Department of Electrical Engineering

University of Cape Town

Submitted to the Department of Electrical Engineering at the University of Cape Town in partial fulfilment of the academic requirements for a Bachelor of Science degree in Electrical and Computer Engineering.

**January 2017**

Blank page

## Declaration

1. I know that plagiarism is wrong. Plagiarism is to use another's work and pretend that it is one's own.
2. I have used the IEEE convention for citation and referencing. Each contribution to, and quotation in, this report from the work(s) of other people has been attributed, and has been cited and referenced.
3. This report is my own work.
4. I have not allowed, and will not allow, anyone to copy my work with the intention of passing it off as their own work or part thereof.

Signature: .....

ROAN SONG

Date: .....

## Acknowledgements

## Abstract

# Table of Contents

<b>Declaration</b>	<b>ii</b>
<b>Acknowledgements</b>	<b>iii</b>
<b>Abstract</b>	<b>iv</b>
<b>Table of Contents</b>	<b>vii</b>
<b>List of Figures</b>	<b>viii</b>
<b>List of Tables</b>	<b>ix</b>
<b>Nomenclature</b>	<b>ix</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Background . . . . .	1
1.2 Motivation . . . . .	2
1.3 Objectives . . . . .	2
1.4 Scope and Limitations . . . . .	3
1.4.1 Focus . . . . .	3
1.4.2 Scope XX . . . . .	3
1.4.3 Limitations XX . . . . .	4
1.5 Report Overview XX . . . . .	4
<b>2 Literature Review</b>	<b>5</b>
2.1 Synthetic Aperture Radar . . . . .	5
2.1.1 Description . . . . .	5
2.1.2 Relevance . . . . .	5

2.2	The MSTAR Dataset . . . . .	6
2.2.1	Description . . . . .	6
2.2.2	Relevance . . . . .	7
2.3	Naïve Classification . . . . .	7
2.3.1	Nearest Neighbour Classification . . . . .	7
2.3.2	K-Nearest Neighbour Classification . . . . .	8
2.4	Deep Learning . . . . .	9
2.4.1	Neural Networks . . . . .	9
2.4.2	Multilayer Perceptron . . . . .	10
2.5	Optimization and Training . . . . .	11
2.5.1	Back-propagation . . . . .	12
2.5.2	Gradient Descent . . . . .	13
2.5.3	Hyper-parameters . . . . .	14
2.5.4	Training, Validation, and Testing . . . . .	14
2.6	Performance Metrics XX . . . . .	16
2.6.1	The Confusion Matrix . . . . .	16
<b>3</b>	<b>Design</b>	<b>17</b>
3.1	Design Context . . . . .	17
3.2	Feasibility Study / Concept Exploration . . . . .	17
3.3	Decomposition and Definition . . . . .	19
3.3.1	Concept of Operations . . . . .	20
3.3.2	User Requirements . . . . .	20
3.3.3	Design Specifications . . . . .	20
3.3.4	High-Level Design . . . . .	22
3.3.5	Detailed Design . . . . .	22
3.3.6	Software Development and Implementation . . . . .	27
3.4	Integration and Recomposition . . . . .	28
3.4.1	Software Development and Implementation . . . . .	28
3.4.2	Subsystem Verification . . . . .	33
3.4.3	System Verification and Deployment . . . . .	33
3.4.4	System Validation . . . . .	34
3.4.5	Operations and Maintenance . . . . .	34
3.4.6	Changes and Upgrades . . . . .	34
3.4.7	Retirement / Replacement . . . . .	35

<b>4</b>	<b>Methodology</b>	<b>36</b>
4.1	Data Processing . . . . .	36
4.2	Dimensionality Reduction . . . . .	37
4.3	K-Nearest Neighbour Classification . . . . .	37
4.3.1	Implementation . . . . .	38
4.4	Multilayer Perceptron . . . . .	39
4.4.1	Implementation . . . . .	39
<b>5</b>	<b>Results</b>	<b>40</b>
5.1	Preliminary Results . . . . .	40
5.1.1	Nearest Neighbour . . . . .	41
5.1.2	K-Nearest Neighbours . . . . .	41
5.1.3	Multilayer Perceptron . . . . .	41
5.2	Preliminary Comments . . . . .	42
5.2.1	Nearest Neighbour . . . . .	42
5.2.2	Multilayer Perceptron . . . . .	42
5.2.3	Data . . . . .	42
5.3	KNN Results . . . . .	43
5.4	Multilayer Perceptron Initial Results . . . . .	43
5.4.1	Analysis . . . . .	45
5.4.2	Performance Metrics . . . . .	45
<b>6</b>	<b>Conclusion</b>	<b>48</b>
	<b>References</b>	<b>49</b>
<b>A</b>	<b>Progress Report</b>	<b>50</b>
A.1	Literature Review . . . . .	50
A.2	Implementation . . . . .	50
A.3	Areas of Focus . . . . .	51



# List of Figures

2.1	Rotational difference between two images of the same class . .	7
3.1	Vee Diagram . . . . .	18
3.2	Nearest Neighbour Classification Flow Diagram . . . . .	24
3.3	K-Nearest Neighbour Classification Flow Diagram . . . . .	25
3.4	K-Nearest Neighbour Optimising for ‘K’ Flow Diagram . . . .	26
3.5	A Neuron in the Network . . . . .	30
3.6	The Internal Workings of a Neuron . . . . .	30
3.7	Multilayer Perceptron Overview . . . . .	31
3.8	Multilayer Perceptron Input Layer . . . . .	32
5.1	Optimising for K . . . . .	44
5.2	Average Cost Over Time . . . . .	46
5.3	Actual Cost Within an Iteration . . . . .	47

# List of Tables

I	Input Data . . . . .	44
II	Initial Confusion Matrix (Before Training) . . . . .	45
III	Transitional Confusion Matrix (Before Final) . . . . .	46
IV	Final Confusion Matrix . . . . .	47

# Chapter 1

## Introduction

This thesis aims to assess the effectiveness of deep learning techniques in the classification of radar imagery. Deep learning relies on the use of neural networks; interconnected layers of nodes sharing information and undergoing non-linear transformations that, through training and optimisation, can detect and extract features from a dataset without human supervision. Training a classifier on a known set of instances allows it to build a predictive model that can then be tested on unseen data, with varying degrees of accuracy. In the specific case of this report, the instances are radar images from the MSTAR dataset.

### 1.1 Background

The desire to mimic human brain function has driven the development of artificial intelligence (AI) and deep learning. The human brain can be viewed as a series of interconnected neurons, firing when undergoing different stimuli. This view led to the foundation of modern deep learning techniques in the 1940s, using multiple layers of artificial neurons. Due to hardware limitations, this approach saw neither success nor widespread adoption. The 21<sup>st</sup> century has seen renewed interest in the field due to increased computational capabilities. While the complexity of the human brain is currently beyond accurate emulation, the deep learning techniques used to

approximate it have found their uses in commercial classification problems.

### 1.2 Motivation

Deep learning is used commercially in voice and image recognition, recommendation engines, artificial intelligence, and a host of other applications. Deep learning classifiers are characterised by relatively long training times and fast classification, making them suited to real-time applications; the time taken to train a predictive model is much greater than the time taken to classify a specific instance, but it can be done beforehand on known data. Acknowledging the success of deep learning techniques has encouraged this report to test the applicability of deep learning techniques in target acquisition and classification of radar imagery.

### 1.3 Objectives

This report aims to develop two classifiers suitable for use on the MSTAR dataset:

1. K-Nearest Neighbours classifier, supporting user-selected values of K
2. Multilayer Perceptron, supporting multiple hidden layers, where each layer's size and activation function can be specified.

Each classifier will be compared according to the following performance metrics:

- Training time
- Training accuracy
- Classification time
- Classification accuracy

The steps to be taken in this report:

- Understand the format of the MSTAR dataset

- Perform pre-processing of the MSTAR dataset
- Develop the KNN classifier
- Develop the Multilayer Perceptron
- Collect performance data for each classifier
- Compare and contrast each classifier based on their performance metrics
- Establish the merit of deep learning techniques in the context of target recognition

## 1.4 Scope and Limitations

### 1.4.1 Focus

This focus of this report is to:

- Review the appropriate academic literature regarding deep learning, and assess the current body of knowledge on the subject to understand where this report will be able to contribute
- Implement KNN and Multilayer Perceptron classifiers in Python
- Implement a training and testing regime for each classifier in Python
- Determine and comment on which classifier is most suitable for the task of radar target recognition

### 1.4.2 Scope XX

Within project scope: Outside project scope:

### **1.4.3 Limitations XX**

All computation will be performed on a desktop computer running Windows 10 Pro with an Intel i5-2500 processor (four cores @ 3.3-3.7GHz), Samsung Evo 850 SSD, and 8GB of DDR3 RAM .

## **1.5 Report Overview XX**

# Chapter 2

## Literature Review

### 2.1 Synthetic Aperture Radar

#### 2.1.1 Description

SAR is used to create images of objects, such as vehicles (as in this report), or landscapes. The images are constructed by sending a radar signal from a moving platform, and the time taken for the signal to return to the antenna denotes the size of the aperture. The aperture can be physical, with a large antenna, or synthetic in the case of a moving aperture. Larger apertures allow for higher image resolution. SAR images consist of magnitude and phase data, from which elevation data can be calculated. [1] The classification of 2D SAR images, the type dealt with in this report, requires only the magnitude data to be preserved.

#### 2.1.2 Relevance

The dataset chosen for this report is comprised of SAR imagery. Understanding the nature of this format allows the decision to strip the data of phase information and keep only magnitude data to be made.

## 2.2 The MSTAR Dataset

### 2.2.1 Description

The MSTAR Public Mixed Targets dataset is provided by the U.S. Airforce on the Sensor Data Management System (SDMS) site [2]. The dataset contains X-band synthetic aperture radar (SAR) image chips of 8 different targets. Each image has a resolution of 1 foot, and is captured in spotlight mode.

The target in each image is centered.

The targets in each class are rotated between  $0^\circ$  and  $360^\circ$ , with images given along the entire path of rotation. This gives a comprehensive view of each target. As Figure 2.1 shows, there is a large disparity in appearance between instances of the same class.

Targets are grouped by elevation angle. For each elevation angle there are between 195 and 274 images per class. Two elevation angles,  $15^\circ$  and  $17^\circ$  were chosen, as they were the two elevation angles for which each class had images (some classes had  $45^\circ$ , but were ignored). Over 8 classes there is a total of 4459 images to consider.

Information pertaining to each target, including its elevation, depression angle, and target type is contained in a header section of each file. The header is followed by magnitude and phase data of the SAR imagery. The SDMS provides tools for converting the raw data into TIFF and JPEG image formats. TIFF is an uncompressed image format, suffering none of the loss that the JPEG format has. It is thus the one used in this report. Converting from the raw data to an image file reduces the complexity of this study to that of image-based target recognition. The phase data present in the original file is safely ignored. [3].

The targets vary in size from  $54 \times 54$  to  $192 \times 193$ , although images within each class are uniform in size.






Figure 2.1: Rotational difference between two images of the same class

### 2.2.2 Relevance

The MSTAR dataset was suggested for this study by A. Mishra. It provides a generic SAR image chip dataset on which any classification method can be run. The dataset is comparatively small, with 195-274 images per class, and thus is suitable for machine learning on a consumer-grade desktop computer with reasonable classification time. The rotation factor present in each image introduces complications, but the images themselves are centered, which is very convenient.

## 2.3 Naïve Classification

### 2.3.1 Nearest Neighbour Classification

#### Description

The nearest neighbour classifier operates as follows: When given an input, the classifier compares this input to the training data set, and finds the one that is closest to the input. For example, if men and women were to be

## CHAPTER 2. LITERATURE REVIEW

classified by their heights, a given input would be classified as either male or female based on the data point in the training data with the height closest to that of the input. This can be expanded to multiple features/dimensions by taking the Euclidean distance between the input and each instance in the training data set. For images, this amounts to comparing, pixel by pixel, each pixel value, and finding the L2 distance between them. Note that there is no need to apply the square root to the distance; it is a monotonic operation, so it will not affect the ordering of the values, and will introduce additional computational complexity. The equation for calculating the L2 distance is:

$$d_2(a, b) = (a - b)^2$$

Each pixel usually has some relation to the pixels near to it, so there is the possibility for a better definition of 'distance' between images to be made [4, 5].

The nearest neighbour classifier has been shown to provide MSTAR classification rates (82-87%) through sufficient image processing and classifier development. SAR images are filled with 'clutter' surrounding the target, which was noted as affecting the success of classification [3].

### Relevance

The nearest neighbour classifier is used in this study as an example of a naïve classifier. Results of this classifier provide a good benchmark against which subsequent classifier performance can easily be measured. The success of other parties in classifying the MSTAR targets using nearest neighbour methods lays a convincing foundation for future development. There is undoubtedly room for improvement, beginning with the elimination of clutter's effect on classification.

### 2.3.2 K-Nearest Neighbour Classification

#### Description

The K-Nearest Neighbour (KNN) classifier is an extension of the Nearest Neighbour classifier where, instead of selecting the single closest neighbour

## CHAPTER 2. LITERATURE REVIEW

from the dataset to the input case, the ‘K’ nearest neighbours are selected, and the class most prevalent amongst the neighbours is taken as the predicted class. K is always an odd-valued integer, to prevent ties from occurring.

### Relevance

KNN is a more robust form of Nearest Neighbour classification because larger values of K ignore outliers in the data. This regularisation effect can be optimised by choosing the best-performing value of K for a specific dataset.

## 2.4 Deep Learning

The objective of this study is to test the performance of deep learning-based classifiers on SAR image chip data. The success of naïve methods has already been proven [3], but lack the predictive power of a more sophisticated classifier. Neural networks and the application of deep learning are key to extracting features from the data to further improve classification rates.

### 2.4.1 Neural Networks

#### Description

A neural network is a system inspired by the perceived workings of the human brain; a system of neurons combine to perform tasks that exceed their individual capabilities. An input is passed through a series of neuron layers, each of which is tuned to identify characteristic features of the input between each layer, allowing for feature extraction and identification.

(From the hartford.edu site) A neural network consists of four main parts [6]:

1. Processing units $\{u_j\}$ , where each  $u_j$  has a certain activation level  $a_j(t)$  at any point in time.
2. Weighted interconnections between the various processing units which determine how the activation of one unit leads to input for another

## CHAPTER 2. LITERATURE REVIEW

unit.

3. An activation rule which acts on the set of input signals at a unit to produce a new output signal, or activation.
4. Optionally, a learning rule that specifies how to adjust the weights for a given input/output pair.

### Relevance

The motivation behind using neural networks is simple; instead of specifying basic characteristics for a system to detect, the system is given an input and a matching output and is left to develop its own perceptions of what important feature link the two. Through optimisation and iteration this can become a very successful form of classification.

### 2.4.2 Multilayer Perceptron

#### Description

A multilayer perceptron is a neural network consisting of an input layer, one or more hidden layers, and an output layer. The input layer is mapped directly from an input instance; one feature per link. In the case of images, each pixel is a feature. The hidden layer neurons have non-linear activation function applied to their inputs, forming their outputs. Typically sigmoid (output range: 0 to 1) or hyperbolic tangent (output range: -1 to 1) are used. The output layer has neurons representing each class and is typically a logistic regression layer; a softmax function is applied to its outputs, making the sum of the layer's neurons' outputs equal to 1. The neuron with the highest output value denotes the predicted class.

Each neuron in every layer is linked to every neuron in the layer that follows it. Each neuron has a randomly chosen weight applied to its output when the classifier is initialised. These weights are subsequently optimised through back-propagation as the classifier is trained. This allows the network to

## CHAPTER 2. LITERATURE REVIEW

develop relationships between neurons, eventually mapping an input to the output of the classifier with as little error as possible.

A single hidden layer is often sufficient for classification tasks, but more layers can be added as desired. Hidden layer 'depth' allows for more complex feature detection, and each layer can be made to serve a purpose as in convolutional neural networks (outside the scope of this report). The multilayer perceptron becomes difficult to optimise as the number of hidden layers grows, because the effect of each neuron on the output, and the effects of previous neurons become progressively more difficult to compute.

### Relevance

Implementing a multilayer perceptron is the main focus of this report, as it provides an example of a deep neural network with a fairly simple implementation. Training time becomes a significant factor when using a deep neural network due to the time taken to complete back-propagation optimisation, so using a multilayer perceptron will likely force the development of more efficient methods of data pre-processing to speed up the training as much as possible.

## 2.5 Optimization and Training

A classifier is only operating efficiently when it is tuned to the data it is attempting to classify. Deep neural networks are initialised with random weights between their neurons, and at first use will perform worse on average than naïve classification methods. Through optimisation of these inter-neuron weights, however, the potential of deep neural networks can be reached, and classification results are expected to significantly improve. Tuning the classifier to the dataset is crucial, but optimising too heavily may result in *overfitting* of the data, leaving the classifier with no predictive power on unseen data.

### 2.5.1 Back-propagation

#### Description

Back-propagation is a system by which the effects of weights between neurons is adjusted through an iterative process. The base case is that of a single input, single output system. Varying the weight on the input directly effects the output. This change can be easily recognised, and the weight can be changed to more suitably link the input to the desired output. This involves developing a method of changing weights in a sensible manner. The most common form of this is through *gradient descent*, covered in Section ??, whereby the weights are adjusted corresponding to their perceived effect on the output state, and their rate of change. Back-propagation is not guaranteed to find a global minimum, and can settle on a local minimum instead, which can be somewhat alleviated through the use of random weights and multiple training rounds, before choosing the best version of the classifier that has been discovered.

One of the key issues with back-propagation is its computational complexity. With deep neural networks, the sheer number of weights and their possible combinations make discerning their impact on the output very difficult, and computationally infeasible to perfectly optimise.

#### Relevance

Back-propagation is a popular and successful technique, well-suited to neural networks with only a few layers. With enough time, it can help to optimise much larger networks, and potentially improve classification accuracy by a large margin. It alleviates the concern of trying to find the perfect network from the offset; it allows any network to be tuned to be better than it currently is.

## 2.5.2 Gradient Descent

### Description

If the computational graph of a classifier’s cost can be calculated, each parameter that contributes to the cost to be adjusted according to its contribution, with the aim of reducing the overall cost. The degree to which each parameter is adjusted is scaled by the *learning rate*. High learning rates lead to rapid change of parameter values, which can result in faster convergence to the optimal values, but can also end up in oscillation around these values if the adjustment is too large. Smaller learning rates can take longer to converge more safely, but can also result in local optimums being converged to, instead of the best possible set of values.

Given a weight  $w$ , a learning rate  $a$ , and the gradient of the cost with respect to the weight,  $\Delta w$ :

$$w = w - a * \Delta w$$

Given a large enough dataset, it may be infeasible to calculate the impact of every weight of every instance on the output, and so the gradient is approximated by taking a batch of instances and averaging the gradient of each, approximating a ‘global gradient’. This is known as *stochastic gradient descent* [7].

### Relevance

Stochastic gradient descent is used in this report in the Multilayer Perceptron to optimise the inter-neuron weight values.

### 2.5.3 Hyper-parameters

#### Description

Hyper-parameters are parameters that, when changed, modify the structure or operation of the neural network, without changing its core mechanics. Hyper-parameters under consideration in this project are:

- K in the KNN (??)
- Input Image size (100px vs 1000px)
- Learning Rate
- Network Shape (hidden layer size, and number of layers)
- L2 regularisation on inter-neuron weights
- Training, Validation and Testing ratios

#### Relevance

Optimising hyper-parameters is incredibly important when optimising a classifier. Sub-optimal hyper-parameters such as a learning rate that is too high could prevent an otherwise functional classifier from converging onto the optimal choice of inter-neuron weights and its classification accuracy could suffer.

### 2.5.4 Training, Validation, and Testing

Successful classification of a dataset is divided into three distinct steps:

1. Training
2. Model Validation
3. Testing

Training is the process of fitting a classifier - it involves running multiple iterations on a given set of inputs, comparing the output of the classifier



## CHAPTER 2. LITERATURE REVIEW

to a known target dataset, and adjusting the parameters of the classifier (typically inter-layer weights and bias) through back-propagation. To find the best iteration of the classifier model, it is periodically tested on a different set of known data; the validation dataset. Testing on this intermediate dataset is used to provide performance metrics such as the mean error and the accuracy of classification, which is useful in selecting a model that is optimised to the desired set of parameters. The data from periodically testing on this validation set is used to tune the model, or implement early-stopping procedures. For example, if the desired level of classification accuracy has been achieved or if the mean error hasn't changed significantly after a number of epochs, the training can be stopped early. On the contrary, if the training is approaching its stated limit yet still improving classification accuracy, the number of epochs or 'patience' can be increased to allow for further iterations and tuning.

Once the training is complete, having achieved the desired level of classification accuracy, the chosen model can be tested on another dataset (typically a set of real-world instances) to see how it performs, providing the testing accuracy. The classifier is no longer tuned, and can be presented with a variety of inputs to simulate its real-world performance.

A dataset is typically split into three sections. 50% training, 25% validation, and 25% testing is a reasonable starting point, and the proportions can be seen as a hyper-parameter to be optimised. If the training set is too small, there is a risk that the model will not be able to successfully extract the features required for classification, and its testing accuracy will be low. If the validation or test sets are too small, the model validation and testing might not be representative of the classifier's performance on a larger test set.

## 2.6 Performance Metrics XX

### 2.6.1 The Confusion Matrix

Confusion matrices are useful tools for evaluating classifier performance. They can show the accuracy, misclassification rate, true/false positive rates, specificity, precision, and prevalence. They clearly show the number of correctly classified classes along the matrix's diagonal, and also shows how many of the incorrect classifications were attributed to which classes. The numbers can be changed to show the percent values of each; i.e. the classification accuracy of each class and how what percent of misclassifications went to each class.

An example confusion matrix is shown in Table II

# Chapter 3

## Design

### 3.1 Design Context

This study is within the scope of an undergraduate level approach to both radar and machine learning, with an emphasis on machine learning; the techniques applied herewith are not limited to radar imagery, but attempts should be made to tailor the design to target recognition applications. With development this study should be adaptable to commercial use, and be helpful to people in the radar department who need assistance with radar target classification. Thus the study should develop an easily extensible framework for radar image classification, or at least guidelines to allow others to integrate with the work covered herein.

### 3.2 Feasibility Study / Concept Exploration

Does the classification of radar imagery lend itself to deep learning techniques, and if so, will the performance be better or worse than naïve classification methods? This is the question that best captures the analysis of the study's feasibility.

Consideration of this question needs to include the format of the data entering the system, the ability of the system to process such data, and

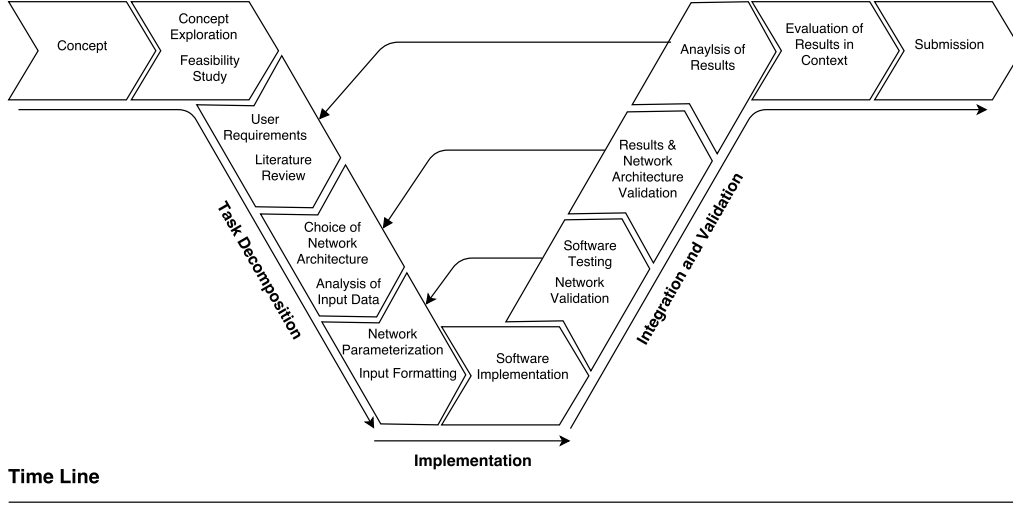


Figure 3.1: Vee Diagram

the effectiveness of the classification of the data.

The data consists of images of between 2916 (54x54) and 37054 (192x193) pixels in size. Each image contains one target, positioned at the centre of the image. This should allow for effective resizing of images to account for size discrepancies. Each pixel in an image constitutes an input. The computational cost of processing the largest image in the set versus the smallest will be at least 12.7 ( $37054/2916$ ) times more expensive. Initial thoughts were that processing of such images will be made much more feasible if they can be reduced in size to match the smallest images present in the dataset, or at least be made as small as possible while retaining all of the information needed to classify each target. However, the shadow of each radar target often extends outside the 54x54 pixel range, yet could be necessary in classification. Preserving features inherent to each image is more important if the classifier can handle the largest images in a timely fashion.

Neural networks have a fixed structure; a collection of input neurons feeding their values through a series of hidden neuron layers before arriving at an output layer of neurons equal in size to the number of classes present in the data. The architecture of the network - the choice of number of hidden layers, the number of neurons in each hidden layer, and the number of neurons in the

input layer are all subject to change during the development of the system. The training and operation of the system occurs through the adjustment of inter-neuron weights, with the structure of the system remaining constant.

The choice of size of the input layer is crucial; it must remain constant throughout the training and operation of the network. Since each pixel in an image forms of the input layer's neurons, all of the input images must be processed to contain the same number of pixels before any other work on the network's architecture can begin.

Once the pre-processing of the input images is complete, the structure of the neural network can be decided. This structure shall be changed and prototyped in order to try to find a good corresponding fit for the data. Having too many hidden layers will greatly increase the time taken by the back-propagation algorithm to optimise the weights of the system, and the likelihood of it settling at a local instead of global minimum increases with the complexity of the system.

The number of neurons in each hidden layer must also be chosen carefully; too many neurons in each layer will result in much longer optimization time (proportional to the increase in the number of inter-neuron weights created). Having too few neurons in a layer can result in the system being unable to extract the features key to classification, and too many neurons may lead to 'overfitting' of the training data, leaving the system with no predictive capability (an inability to classify data not present in the set of training instances).

### 3.3 Decomposition and Definition

This section is devoted to describing the study in terms of its requirements, operation, and implementation. An accompaniment explaining the verification and validity of each subsection will be in the next section.

### 3.3.1 Concept of Operations

Radar target classification is an inexact science; interpreting a radar image and comparing it to a known case is not as straightforward in all cases as one might expect. Weather conditions, environmental clutter, and image resolution all obscure the target to varying degrees, making intuitive classification ineffective. Computer-based classification through analysis of multiple targets and the application of deep learning techniques should in theory allow distorted images to be classified after the computer is trained to recognise features pertaining to each class. Naïve methods of classification lack predictive power - the ability to 'guess' effectively if the target is obscured or unrecognised. Deep learning methods are the solution that this study proposes.

### 3.3.2 User Requirements

The success of this report is based on the ability to correctly classify and recognise radar targets taken from the supplied dataset using deep learning techniques.

The following is required:

- Use the MSTAR dataset (Section 2.2)
- Develop a naïve classifier to use as a benchmark
- Use deep learning methodology to develop a classifier
- Indicate how the classifiers can be improved
- Comment on the performance of each classifier
- Report on the suitability of deep learning for target recognition

### 3.3.3 Design Specifications

Expanding upon the user requirements, the following design specifications have been derived:

**The system must be trained on the MSTAR database of radar images**

The images in the MSTAR dataset have undergone a level of pre-processing, making them suitable for rapid prototyping and development of classifiers.

**The system must have a testing accuracy of above 95%**

Correct identification of radar targets is the aim of this entire report, and as such is the most important performance metric to consider.

**The Nearest Neighbour classifier should be used as a benchmark for classifier comparison**

The NN classifier provides a good example of a naïve classifier, and should be beaten by any classifier that is somewhat optimised for the dataset in use. The NN classifier serves as a good benchmark to improve upon, as it gives a lower bound of expected results

**At least two different classifiers should be tested against the Nearest Neighbour classifier**

The chosen classifiers are the K-Nearest Neighbour classifier optimised to fit the dataset, and the Multilayer Perceptron, which satisfies the need for a deep learning classifier.

**Each classifier must be evaluated and compared**

The performance metrics to be used are training time, classification time, and classification accuracy.

### 3.3.4 High-Level Design

After analysis of the user requirements, the following areas of design need to be focused on:

- Identification/classification
- Image Processing/Preparation
- Dimensionality Reduction
- Naïve Classification (Nearest Neighbour as a benchmark)
- Deep Learning Classification (Multilayer Perceptron)
- Obtaining classifier performance metrics
- Classifier comparison
- Classifier optimisation

### 3.3.5 Detailed Design

#### Image Processing/Preparation

The MSTAR dataset is a compilation of image chips, all of which contain a header, as well as magnitude and phase data. The images are between 54x54 and 192x193 pixels in size, which suggests that some form of image processing should be performed to make sure that all images are the same size. The targets in each image chip are centred, suggesting that cropping each image to a size where the target (and its shadow - useful in classification) are left whole, and as much of the surrounding clutter as possible is removed.

An alternate approach is to retain the data inherent in the environmental clutter and pad the smaller images with zeros, keeping all images in the set at the size of the largest image in the set. While this preserves all of the image chip data, processing larger images leads to longer training and classification times.

A compromise is to ensure that all images are the size of the largest image in the set, and somehow reduce the clutter present in each image. Because the



targets are substantially brighter than their surroundings, using some form of thresholding (setting values lower than a specified threshold to zero) should prove to be effective in lessening the impact of the clutter, if not completely removing it.

### **Dimensionality Reduction**

Each pixel in an image is taken as an feature, forming a feature vector with a length equal to the total number of pixels in the image. The image cropping mentioned in Section 3.3.5 is very effective at reducing the size of this feature vector. If a 128x128 image is cropped to 64x64, the feature vector's length is reduced by factor of 4. This can be reduced further through the application of dimensionality reduction techniques, such as Principal Component Analysis, Locally-Linear Embedding, Sum of Means and non-linear methods, all of which are outside the scope of this report. As such, each image will be rescaled and padded with zeros to match the size of the largest image in the MSTAR dataset (192x193). This allows every possible feature in each image to be used, while allowing for dimensionality reduction if desired.

### **Nearest Neighbour Design**

The high-level design is shown in Figure 3.2. Its principles of operation are covered in Section 2.3.1. To implement this classifier, the following is needed:

- Access to the dataset
- A choice of input image
- A method to calculate and sum the pixel-wise distances
- A variable storing the smallest distance and tentative classification
- A method displaying the chosen class and whether or not it is correct




Figure 3.2: Nearest Neighbour Classification Flow Diagram

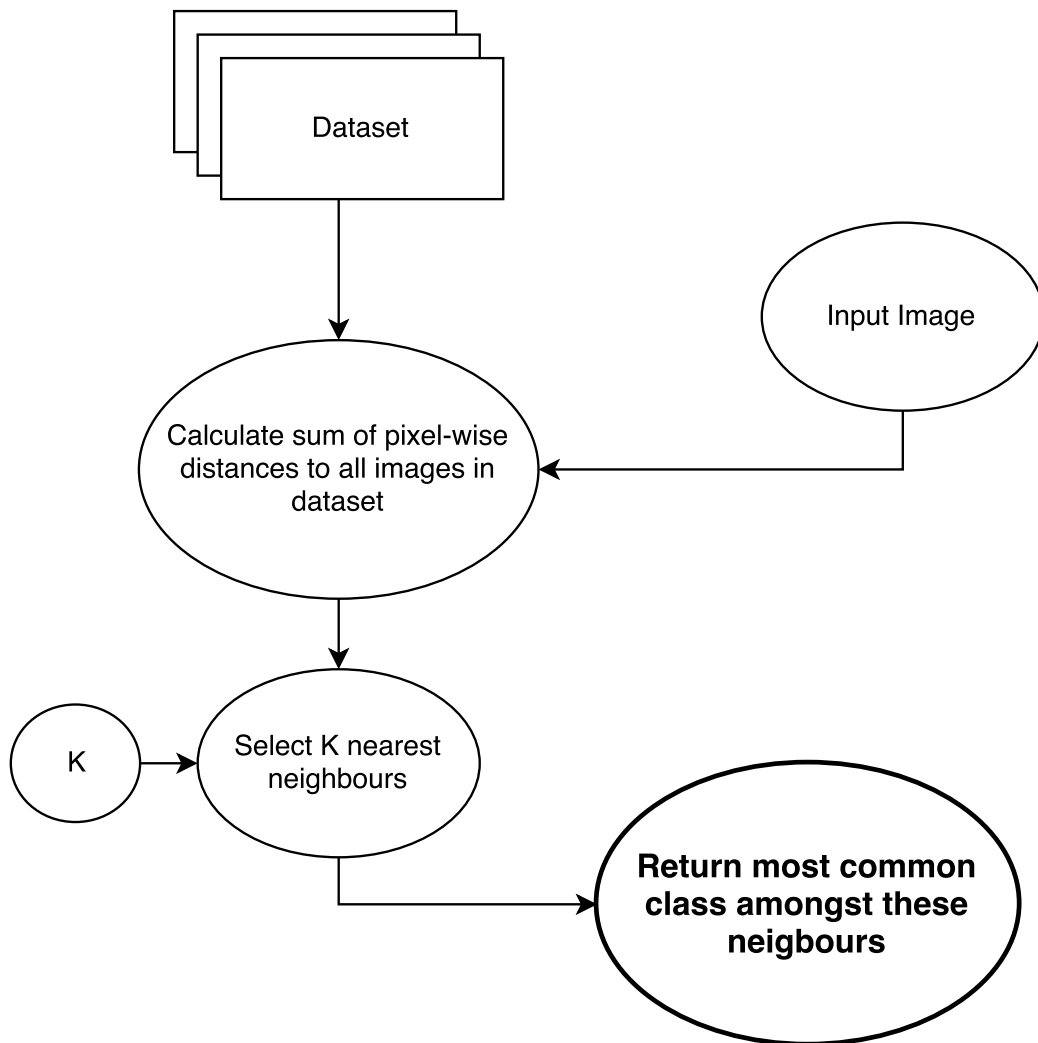


Figure 3.3: K-Nearest Neighbour Classification Flow Diagram




Figure 3.4: K-Nearest Neighbour Optimising for 'K' Flow Diagram

### **K-Nearest Neighbour Design**

While similar to the Nearest Neighbour design, KNN introduces its own complexities, most notably when optimising for K. The high-level design is shown in Figure 3.3, its principles of operation are covered in Section 2.3.2. The method of finding the best K value for a given dataset is shown in Figure 3.4.

### **Multilayer Perceptron Design**

Implementation of a multilayer perceptron in software can be divided into discrete sections as follows:

- Implement a wrapper class for the classifier
- Develop a logistic regression layer class to use at the output
- Develop an extensible hidden layer class
- Collate the important details of each layer in the wrapper class
- Implement a test method

### **3.3.6 Software Development and Implementation**

Developing the software for this report is done with the following objectives in mind:

- Comment code clearly
- Use logical code structure and layout
- Use Python as the chosen programming language
- Use Theano module for deep learning calculations
- Code self-contained classifier classes
- Difficult to generate data should be saved for future use
- Code reusable, adaptable methods

- Make performance metrics readily available
- Test while developing
- Provide a comprehensive ‘README’ for new users

The software implementation will be covered in Section 4. The software present in this study has been made available on Github at [github.com/roansong](https://github.com/roansong).

## 3.4 Integration and Recomposition

### 3.4.1 Software Development and Implementation

#### Inline Testing

While implementing the software for this report, I ran into some early issues that resulted from insufficient planning. After taking a step back, it was decided that more steady progress would be made by implementing rigorous inline testing, i.e. incrementally testing the code after every slight modification, instead of only testing after the addition of a major feature and then trying to iron out any latent bugs. This approach leads to much simpler debugging.

#### The Multilayer Perceptron

In a process outlined in Figure 3.8, the input image is converted to an array of unsigned 8-bit integers (ranging from 0 to 255). The elements of the array are ‘unitised’ by subtracting the mean of the array from each and then dividing each element by the standard deviation of the array. This ensures that the processed values lie centred around zero, and mostly between -1 and 1. Unsigned integers are no longer suited to represent this data, so 32-bit floating point numbers (floats) are used.

Once each pixel has been unitised, they form the input layer of the neural network, with each pixel representing a neuron in the input layer shown in

Figure 3.7.

Each of these input layer neurons has a random weight applied, and is then fed into the neurons that form the first hidden layer. Every neuron in the input layer contributes to every neuron in the next layer. At each neuron in the first hidden layer, the ‘net’ value is formed by summing all of the values present at its input (i.e. all the weighted values passed along from the input layer). An activation function is applied to this net value to compress the range of values. There are many different activation functions that can be used, and for this study I narrowed down the choice to either the logistic function  $\frac{1}{1+e^{-net}}$  or the hyperbolic tangent function  $\frac{e^{net}-e^{-net}}{e^{net}+e^{-net}}$ . The logistic function produces an output between 0 and 1, while the hyperbolic tangent function’s output is between -1 and 1. I decided to use the hyperbolic tangent function, because the logistic function has a tendency to incur long training times if its values lie very close to 0, while the hyperbolic tangent function tends to move towards its extreme values more quickly.

The activation function, when applied to the net of the neuron, becomes the output of the neuron, feeding through to every neuron in the next layer with weights applied, repeating the same process until the output layer. This is shown in Figure 3.5 and more closely in Figure 3.6.

The output layer has as many neurons as there are classes. The output layer is typically a logistic regression layer, which has a softmax function applied to its outputs. This results in outputs that sum to 1. The neuron with the highest value is taken to be the network’s prediction. An example output would look something like:

$$[0.25, 0.2, 0.1, 0.05, 0.4]$$

There are five classes, and the fifth class has the highest value, thus the neural network has classified the input as belonging to that fifth class. Given that the maximum possible value for an output would be 1, 0.9 would show a large confidence in the classification. A value close to 0 shows strong disagreement, and something like 0.3 would show moderately low confidence.

To train the network, stochastic gradient descent is used, as described in Section ??.




Figure 3.5: A Neuron in the Network




Figure 3.6: The Internal Workings of a Neuron

### Algorithm Testing

For each classifier, their effectiveness on the MSTAR dataset must be tested and catalogued, beginning with the Nearest Neighbour implementation. Doing the Nearest Neighbour classification first establishes a benchmark against which further methods can be tested. The naïve nature of the Nearest Neighbour classifier means that it is not optimised for any particular dataset. It shows the efficacy of a generic algorithm applied to the MSTAR dataset.

For the effectiveness of any algorithm to be tested, correct and incorrect outputs must be defined. The MSTAR dataset includes target labels



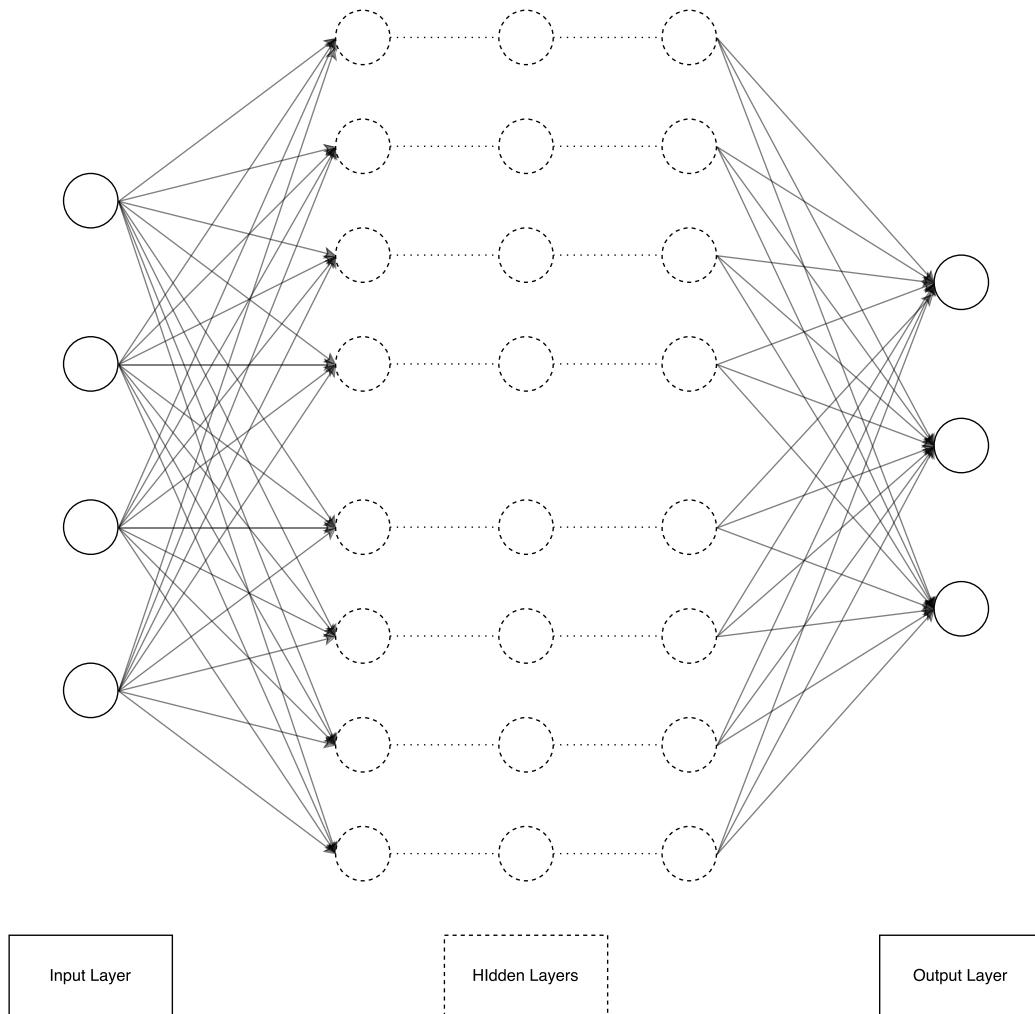


Figure 3.7: Multilayer Perceptron Overview

in a header section of each file, but since the operations are conducted on files with their header and phase data stripped away it merely adds computational complexity to find the corresponding header for each file and then parse it to extract data about each target. To simplify classification, a variant of a one-hot vector denoting the classes is attached to each target. The vector consists of a series of numbers, equal in length to the number of classes in the dataset. A ‘1’ denotes that the instance is a member of the class corresponding to that entry in the vector, and the rest of the numbers are 0, showing that the instance is not in those classes. A file is created listing all of the filenames to be tested during the run of the algorithm. An

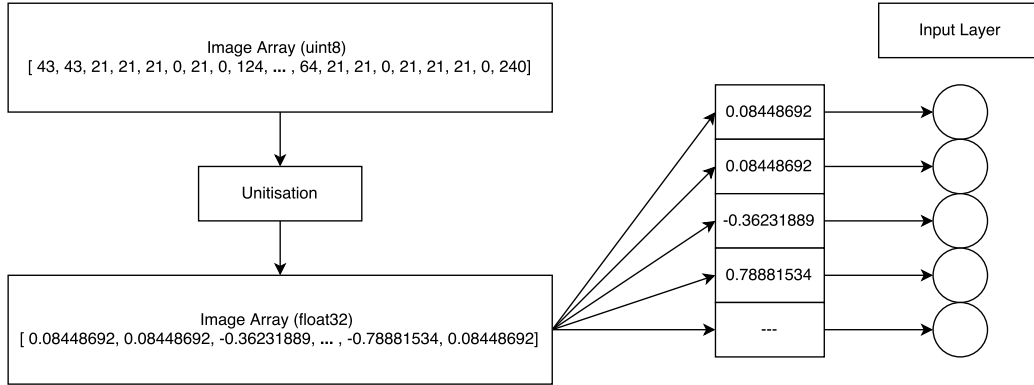


Figure 3.8: Multilayer Perceptron Input Layer

example file with ten entries and two classes would look as follows:

```
HB03333.003.tiff 1 0
HB03334.003.tiff 1 0
HB03335.003.tiff 1 0
HB03337.003.tiff 1 0
HB03338.003.tiff 1 0
HB14931.025.tiff 0 1
HB14932.025.tiff 0 1
HB14933.025.tiff 0 1
HB14934.025.tiff 0 1
HB14935.025.tiff 0 1
```

### Testing During Development

The simplest way to find a classifier's efficacy is to test it on a wide variety of classes and on as many test instances as possible. To provide interim results, during the iterative phase of classifier development, only a subset of the dataset's images are used. This compromises the final accuracy of the classifier (it may perform differently on the full dataset), but brings with it the ability to test and train classifiers more quickly, due to the lower computational overheads. During the development this testing method allows for simple decisions regarding the direction of classifier implementation

or optimisation to be made. In process of verifying the classifiers, the full dataset must be used to give an accurate picture of the classifier's performance and real-world implementation.

### 3.4.2 Subsystem Verification

The chosen method for verifying a classifier's efficacy once it is considered to be sufficiently optimised is simple; The classifier is tested on the dataset with a number of test instances as in Section 2.5.4. For the KNN, fully testing it entails removing one instance from the dataset, training the classifier on the remaining points, and using the removed instance as a test case. Once this has been done, the instance is replaced, another is taken, and the process is repeated until every instance in the dataset has been tested. The system's performance is based on the number of correct classifications made during the process. This is known as Leave-One-Out Cross Validation (LOOCV).

### 3.4.3 System Verification and Deployment

To confirm the results of each classifier, it is important to have a set of training instances on which the system can be trained. The system is then tested on another set of points whose classes are known. Once this has been tested and confirmed to have a desirable level of classification accuracy, the system will be ready for testing on previously unseen, real-world cases.

This is accomplished by dividing a set of known data points into a training set, a validation set, and a testing set. For example: 80% of the data points will be used to train the system, 10% to validate the model, which will then be tested on the remaining 10%. The process of cross-validation entails selecting a different training/test split each time (either systematically or at random) and performing the process again. This concept can be extended to where the system is trained on all but one of the instances and then tested against it, which is known as LOOCV ("Leave One Out" Cross-Validation). The system is tweaked until it reaches the level of classification required. Cross-Validation is an important tool for eliminating "overfitting" of the system to the training data. Mixing up the training and test cases ensures that the

classifier is left with some ability to generalise, and not just repeat what it has been shown.

### 3.4.4 System Validation

If the objectives of the report have been accomplished, the system will be considered valid. The results found during this report have to be collated and analysed within the context of its requirements. This goes beyond confirming the individual results of each classifier and seeing if implementing deep learning techniques are indeed the ‘smart choice’ in the task of radar target recognition.

### 3.4.5 Operations and Maintenance

Considering the real-world application of a target recognition system, it should meet certain criteria to ensure ease of use and compatibility with various datasets. Providing adequate documentation to support the system is key.

The documentation must detail:

- The operating procedure of the system
- The expected input to the system (dataset and individual instance)
- The output format of the system (predicted class and performance metrics)
- Comprehensive troubleshooting (outside the scope of this report)

### 3.4.6 Changes and Upgrades

If the system is required to be maintained and expand its scope (by incorporating more, larger, and more complex images into its dataset), the input data will have to undergo pre-processing that is not currently implemented into the system, such as intelligent dimensionality reduction.

The system may have to restructure some of its key classification methods and re-optimize hyper-parameters.

A classification method to consider for real-world use is a Convolutional Neural Network, as it can classify targets without the need for the target to be centered (as in the MSTAR dataset). This would be more suited to real-world applications, although the classification time is longer. This is not trivial, and lies outside the scope of this report.

### **3.4.7 Retirement / Replacement**

# Chapter 4

## Methodology

### 4.1 Data Processing

The first step is to get to grips with the data - processing the data set in a way that makes sense to use, and allows different classification methods to be implemented easily on it.

The MSTAR dataset contains eight different targets. The dataset is sorted by depression angle, and by ‘scene’. All of the targets have data corresponding to 15° and 17° elevation angles, for a total of 4459 images.

For unsupervised learning, we want to have each target in the dataset labelled with corresponding information, most notably its class. This is done using a one-hot array (i.e.  $[1 \ 0 \ 0]$ ,  $[0 \ 1 \ 0]$  for 1 and 2, respectively) relating to each target. The MSTAR dataset stores the information of each target in a header section of each file. This is inconvenient when reading in image files directly, so the header is discarded and images are classed according to their file extension. The file extensions and the classes they correspond to can be found in Table I.

## 4.2 Dimensionality Reduction

An image can be viewed as a collection of pixels, comprising a feature vector. If the size of this vector can be reduced while retaining enough information for classification, we can greatly increase our training and classification speeds. Simply selecting every other pixel (or one in five) would greatly reduce the dimensionality of the data but may remove features that are key to classification, thus having a negative impact on the results.

## 4.3 K-Nearest Neighbour Classification

The steps necessary to implement the NN on the MSTAR dataset are as follows:

- Convert the raw data + header MSTAR files into .tiff images
- figure out how to read and display these images in matlab
- write a pixel-by-pixel comparison method
- test the method by comparing one image to itself and others
- collate a set of mixed radar targets for testing
- extract data from each image/filename to help with seeing if the classification is correct
- TEST
- collect results!

The KNN classifier works by predicting the class of an instance based on its relative distance to nearby instances. The K nearest instances (the “neighbours” after which the classifier is named) are used, and the most common class amongst them is selected as the predicted class. K is always an odd number, to prevent ties. A measure of confidence in the prediction can be taken as the ratio of the most occurring class to the total number of neighbours under consideration. K can be optimised for each dataset

### 4.3.1 Implementation

#### The algorithm

Given a training dataset of known instances and targets, and an input instance  $X$ , the squared difference between  $X$  and each instance in the dataset is calculated. The  $K$  closest instances (neighbours) are selected, and the class predicted is the most commonly occurring class within the selection of neighbours. The time taken to predict a class is linearly proportional to the size of the training dataset ( $O(n)$ ), and exponentially proportional to the size of each instance ( $O(n^2)$ ). A 100x100 image has 10,000 points of comparison; 100 times more than a 10x10px image. Reducing the image size through cropping or dimensionality reduction can thus have a desirable effect on computation time.

#### Optimising for $K$

Optimal  $K$  values are found by performing leave-one-out cross-validation (LOOCV) and finding which value of  $k$  gives the greatest prediction accuracy across all cases. To find this value of  $k$  using the previous algorithm is a very time-consuming process. Because this optimisation uses only the training set, the process can be sped up significantly. A matrix containing the squared distances between each instance is constructed. An example of such a matrix is shown in Figure ???. Each row represents the squared distances between the instance corresponding to that row, and every other instance. This leads to the matrix being mirrored along its leading diagonal, which speeds up the construction of the matrix by not having to recalculate those values. All of the diagonals of the matrix are zero; they represent the distance between an instance and itself. To optimise for  $K$ , each row is taken in turn, sorted in ascending order, removing the first element (the zero self-term), and calculating the predicted class and confidence.



## 4.4 Multilayer Perceptron

### 4.4.1 Implementation

#### The wrapper class

The wrapper class encapsulates the layers of the Multilayer Perceptron, specifying the the dataset to be used, the size of each layer, and the number of classes.

#### The logistic regression layer class

The logistic regression layer serves as the output layer for the Multilayer Perceptron, providing a convenient way to see the predicted output that corresponds to a given input, and to test the error of this output, given by the negative log likelihood between the predicted output and the target output for the input image.

#### The hidden layer class

Hidden layers are the fundamental building block for deep neural networks. One or more can be used, connected in series from the input to the output, with each neuron in a layer connecting to every neuron in the next. Each neuron in a hidden layer applies a non-linear activation function to the sum of its inputs to produce an output

- Develop an extensible hidden layer class that will be between the input and output, and can connect to more hidden layers en route.
- Collate the important details of each layer (such as the logistic regression layer's predicted class) in the wrapper class for ease of access
- Implement a test method with variable parameters for regularisation

# Chapter 5

## Results

This section of the report shows the results of the tests performed on the chosen classifiers

### 5.1 Preliminary Results

The preliminary results were calculated on a subset of the MSTAR dataset; 1291 images spanning 5 classes were selected. Each image was cropped/padded to 100x100 pixels in size and then normalised, forming an input of 10,000 data points. No further processing was performed on the images.

The data was divided as follows: 50% training, 25% validation, 25% testing. Instances are drawn randomly from the dataset without replacement.

### 5.1.1 Nearest Neighbour

### 5.1.2 K-Nearest Neighbours

### 5.1.3 Multilayer Perceptron

The multilayer perceptron has the largest number of parameters that can be optimised, leading to a varied spread of results. The first parameters to optimise are the learning rate and the number of neurons in the hidden layer(s).

The learning rate controls how strongly the output error affects the shift in weight values every iteration. A smaller learning rate will typically cause the network to take longer to converge to its optimal weights, or get stuck in a local minimum, while a high learning rate can prevent the values from converging, resulting in oscillating values and no optimum being found.

The size of each hidden layer defines the number of connections between neurons. Too few connections can ‘bottleneck’ the classifier, leaving it unable to extract any meaningful features. Too many connections can allow the classifier to extract more features, at the cost of a longer training time spent optimising the larger number of weights. It is predicted that too many connections can also lead to over-fitting of the classifier to the training dataset.

Preliminary results show that learning rate must be set in indirect proportion to the number of hidden neurons to maintain training and classification accuracy.

**Single Hidden Layer**

**Multiple Hidden Layers**

## **5.2 Preliminary Comments**

Analysis of the preliminary results gives direction to the development of each classifier.

### **5.2.1 Nearest Neighbour**

The Nearest Neighbour classifier results are intended to function only as a low benchmark and point of reference, so its design remains unchanged.

### **5.2.2 Multilayer Perceptron**

The required hidden layer size is much smaller than anticipated; a single hidden layer with 10 neurons performs adequately. Optimisation of the weights according to their L2 norm proved crucial in improving predictive accuracy while reducing the cost of the classifier.

### **5.2.3 Data**

The division of the dataset between training, validation and testing should be changed to 85% training, 5% validation, 10% testing. This is based on the idea that more training data results in a stronger classifier. The validation stage is for verifying the progress of the model and does not need too much data dedicated to it. Testing requires more data than validation - enough to be confidently representative of the contents of the entire dataset.

### 5.3 KNN Results

The KNN algorithm is simple to execute, but scales poorly with dataset size and image size. Using a dataset of 1291 images, of size 100x100, the largest bottleneck in the procedure is the generation of the squared distance matrix  $D_2$ , taking 182 minutes (3 hours and 2 minutes) to complete. This was then saved to prevent the need to recalculate it in future. The time taken to calculate the optimal value of  $K$  is also proportional to the number of images, taking 379 minutes (6 hours and 19 minutes) when considering all odd values of  $K$  up to and including 1291. Plotting these  $K$ -values vs the training error obtained for each value is shown in Figure 5.1. The time taken to classify a single instance is 0.25m (15s). Testing though the full set of 1291 to obtain the training classification accuracy takes  $\sim 1291 \times 0.25$  minutes, resulting in a real-world time of 342 minutes (5 hours and 42 minutes).

### 5.4 Multilayer Perceptron Initial Results

The results in this section use an input image of 100x100px. The image is cropped around the center of the image. For some images smaller than this size (e.g. the SLICY set at 54x54px) the images are padded with zeroes to ensure a consistent input image size. The multilayer perceptron has two hidden layers, of 30 and 20 neurons, respectively. The resulting structure is thus 10,000 input neurons feeding into 30, feeding into 20, feeding into the five output neurons. Back-propagation through all layers is applied after every forward run, which has effects on the output that can be seen in 5.4.1. The cost function is the square of the error. The system is trained for 10,000 iterations, and settles after approximately 300. The input images are taken in order (i.e. 195 BTR 60 images, then 274 2S1 images, etc.) per iteration. The layout of the input data is shown in Table I.

## CHAPTER 5. RESULTS




Figure 5.1: Optimising for K

Class	No. Images	Suffix
2S1	573	000
BRDM	572	001
BTR 60	451	003
D7	573	005
SLICY	572	015
T62	572	016
ZIL131	573	025
ZSU_23_4	573	026

Table I: Input Data

### 5.4.1 Analysis

The final values of the system lie at 90.6% classification accuracy (1170/1291 correct). The final confusion matrix (Table IV) shows a clear pattern; the incorrect predictions are always to the class preceding the correct class. Figure 5.2 shows the average cost decreasing over time, while Figure 5.3 shows how the actual cost sharply increases at the start of a new class and decreases swiftly, before rising again at the beginning of the next class. This appears to be due to the layout of the input data. For subsequent development I need to investigate using batches of inputs to reduce this. Reducing the number of consecutive inputs of the same class should force the network to detect underlying features and mitigate this error.

### 5.4.2 Performance Metrics

		Predicted Class				
		BTR 60	2S1	BRDM 2	D7	SLICY
Actual Class	BTR 60	0	195	0	0	0
	2S1	0	274	0	0	0
	BRDM 2	0	274	0	0	0
	D7	0	274	0	0	0
	SLICY	0	274	0	0	0

Table II: Initial Confusion Matrix (Before Training)

## CHAPTER 5. RESULTS




Figure 5.2: Average Cost Over Time

		Predicted Class				
		BTR 60	2S1	BRDM 2	D7	SLICY
Actual Class	BTR 60	170	0	0	0	25
	2S1	23	251	0	0	0
	BRDM 2	0	24	250	0	0
	D7	0	0	24	250	0
	SLICY	0	0	0	24	250

Table III: Transitional Confusion Matrix (Before Final)






Figure 5.3: Actual Cost Within an Iteration

		Predicted Class				
		BTR 60	2S1	BRDM 2	D7	SLICY
Actual Class	BTR 60	170	0	0	0	25
	2S1	24	250	0	0	0
	BRDM 2	0	24	250	0	0
	D7	0	0	24	250	0
	SLICY	0	0	0	24	250

Table IV: Final Confusion Matrix

## Chapter 6

## Conclusion

# References

- [1] J. C. Curlander and R. N. McDonough, *Synthetic aperture radar*. John Wiley & Sons New York, NY, USA, 1991.
- [2] U. A. Force, “The mstar dataset,” 2015. [Online]. Available: <https://www.sdms.afrl.af.mil/index.php?collection=mstar>
- [3] R. Schumacher and K. Rosenbach, “Atr of battlefield targets by sar classification results using the public mstar dataset compared with a dataset by qinetiq, uk summary.”
- [4] L. Wang, Y. Zhang, and J. Feng, “On the euclidean distance of images,” Ph.D. dissertation, School of Electronics Engineering and Computer Science, Peking University, 2005. [Online]. Available: <http://www.cis.pku.edu.cn/faculty/vision/wangliwei/pdf/IMED.pdf>
- [5] D. Michie, D. J. Spiegelhalter, and C. Taylor, “Machine learning, neural and statistical classification,” 1994.
- [6] I. Russel, “Definition of a neural network,” 1996. [Online]. Available: <http://uhaweb.hartford.edu/compsci/neural-networks-definition.html>
- [7] L. Bottou, “Large-scale machine learning with stochastic gradient descent,” in *Proceedings of COMPSTAT’2010*. Springer, 2010, pp. 177–186.

# Appendix A

## Progress Report

As of the 16th of September, certain progress has been made in three areas of the report, namely the Literature Review, Methodology, and Implementation.

### A.1 Literature Review

The Literature Review has been somewhat written; the core topics of most relevance to the study have been covered, but the level of detail for each topic varies, and I feel it may be insufficient. This is a section that I feel needs further refinement and content, that will be added as the development of this study touches on more topics.

### A.2 Implementation

I have implemented a working nearest neighbour classification system, but it is quite slow. I have implemented the first few stages of a multilayer perceptron with two hidden layers, and am currently busy with the implementation of a back-propagation algorithm to optimise the network.

### A.3 Areas of Focus

The next avenue of development is to process the MSTAR image chip data to facilitate quicker classification. Until now I have been using only a subset of the MSTAR data (70 image chips), but this is not feasible in the final product, as it does not include sufficient training data. The time taken to process this subset is non-negligible, and something must be done to improve the training time. The idea is to process the images in such a way that the resolution of each image, and therefore the number of pixels to process, can be reduced while retaining all the crucial data contained in each image. This will include removing the clutter surrounding each target, and may involve clustering algorithms to preserve important areas of data.

I also need to finish my implementation of the multilayer perceptron, specifically the application of the back-propagation algorithm to optimise the network.